

Moral Susceptibility and Robustness in Large Language Models

Davi Bastos Costa, Felipe Alves & Renato Vicente
TELUS Digital Research Hub
Center for Artificial Intelligence and Machine Learning
Institute of Mathematics, Statistics and Computer Science
University of São Paulo
{davi.costa, falves, rvicente}@usp.br

October 27, 2025

ABSTRACT

We study how persona conditioning influences the moral judgments produced by large language models (LLMs). Using the 30-item Moral Foundations Questionnaire (MFQ-30), we elicit repeated ratings across diverse personas and models, and introduce a benchmark that quantifies two properties: (i) moral susceptibility (the extent to which MFQ subscale scores shift under different personas), and (ii) robustness (the stability of ratings under repeated sampling and persona variation). We describe a simple, reproducible experimental protocol and propose variance- and effect-size-based metrics alongside mixed-effects analyses to isolate persona-related variance components. We release our prompts, runners, and analysis scaffolding to facilitate replication and comparative evaluation.

group/Loyalty, Authority/Respect, and Purity/Sanctity (Graham et al., 2009; Haidt & Graham, 2007). We operationalize moral susceptibility as the variation in MFQ subscale scores across personas, and robustness as the stability of ratings across repeated trials and persona perturbations. Our contributions are:

1. A standardized, open protocol for eliciting MFQ-30 ratings from LLMs under persona conditioning, including prompts and a lightweight runner.
2. A set of susceptibility and robustness metrics grounded in variance components, effect sizes, and reliability analysis.
3. An empirical study across multiple models and personas, with guidance for statistical analysis and reporting.

1 INTRODUCTION

Reliable benchmarks for the social capabilities of large language models (LLMs) are increasingly important as these systems are deployed in interactive, multi-agent settings where outcomes hinge on social intelligence and strategic reasoning. Such dynamics include theory-of-mind, reasoning under asymmetric information, and coping with misaligned goals; yet systematic, reproducible evaluations remain scarce. Motivated by this need—and echoing calls to rigorously benchmark social behavior in LLMs (Costa & Vicente, 2025)—we focus on moral judgment as a core facet of social decision-making and alignment.

This paper introduces a benchmark based on the Moral Foundations Questionnaire (MFQ-30), a widely used instrument in moral psychology that measures five moral foundations: Harm/Care, Fairness/Reciprocity, In-

Recent MFQ-based studies profile LLM value orientations and alignment. Abdulhai et al. (2024) adapt MFQ prompts to derive foundation scores, compare them to human surveys, and show that targeted prompts can shift profiles and affect downstream donations. Nunes et al. (2024) combine MFQ with MFV to reveal inconsistencies between abstract and concrete judgments. Aksoy (2024) use MFQ-2 across eight languages to expose cultural/linguistic variability, and Bajpai et al. (2024) compare MFQ-20 and moral competence between humans and chatbots, finding LLMs emphasize individualist foundations and lag human competence. In parallel, MoralBench (Ji et al., 2025) offers a broad task suite; our MFQ persona framework complements it by isolating persona-driven shifts relative to a self baseline.

2 MORAL ROBUSTNESS AND SUSCEPTIBILITY BENCHMARK

We define a benchmark to evaluate two complementary dimensions of persona sensitivity in LLMs.

Robustness The stability of MFQ ratings under repeated sampling and small persona perturbations (e.g., paraphrases). Operationally, we report a simple index defined as the inverse of the average per-item standard deviation across repetitions (higher is more stable).

Moral susceptibility The degree to which MFQ subscale scores shift as persona descriptions change. High susceptibility indicates strong persona-driven modulation of moral judgments; low susceptibility indicates persona-invariant responses.

2.1 MFQ

The MFQ-30 comprises 30 items split into two sections: 15 relevance judgments (how relevant specific considerations are when deciding right from wrong) and 15 agreement statements (level of agreement with moral propositions) (Graham et al., 2011). Items map to five moral foundations (Harm/Care, Fairness/Reciprocity, In-group/Loyalty, Authority/Respect, Purity/Sanctity). Following common practice, filler items (e.g., canonical item indices 6 and 22 in some MFQ-30 versions) are excluded from subscale scoring. Subscale scores are computed by averaging the items associated with each foundation within each section and then combining sections (e.g., mean of relevance and agreement for that foundation), or by an alternative pre-registered scheme.

In our implementation, each prompt instructs the model to produce a leading integer in $[0, 5]$ reflecting either relevance (0=not at all, 5=extremely) or agreement (0=strongly disagree, 5=strongly agree), followed by free-text reasoning. Ratings are parsed by extracting the first digit $[0, 5]$ from the response. Figure 1 illustrates the resulting MFQ relevance profile across models using the self (no-persona) baseline.

2.2 Experimental Methodology

We use a simple, reproducible runner that iterates through MFQ-30 items for a list of personas and repeats each item multiple times to characterize response variability. The runner supports local GGUF models as well as API-hosted models through a uniform interface. Concretely:

- **Personas:** A JSON file provides persona descriptions (plain text). By default, each persona is used as-is and identified by its index.




Figure 1. Moral foundation relevance profiles (self/no-persona baseline). Points show mean relevance per foundation; error bars denote standard errors across items within each foundation.




Figure 2. Moral foundation relevance profiles for five randomly selected personas, averaged across models. This visualization highlights an averaged effect of persona identity (persona_id) on MFQ relevance patterns. See the Personas appendix for descriptions; indices match the zero-based persona_id used in our runs.

- **Prompting:** For each persona and item, the model receives a roleplaying instruction (“You are roleplaying as the following persona: ...”) plus the MFQ item prompt. The prompt requests a leading integer rating in $[0, 5]$ and then reasoning.
- **Repetitions:** Each persona–question pair is queried n times (default $n = 10$) to estimate within-persona variability and enable reliability analysis.
- **Decoding:** We use low temperature (default 0.1) and a small `max_tokens` (default 5) to elicit short, rating-first outputs. Ratings are parsed with a conservative regex; failures are recorded as -1 .
- **Logging:** Each response is streamed to CSV with fields: `persona_id`, `question_id`, `run_index`, `rating`, `truncated response text`, and `timestamp`.
- **Models:** We include local chat-tuned GGUF models (e.g., Mistral, Llama, Qwen) and hosted models (e.g., Anthropic, OpenAI) when API keys are configured.

2.3 Statistical Analysis

This section formalizes the quantities we compute from the MFQ runs and how we summarize them into moral robustness susceptibility metrics with uncertainty.

Let \mathcal{P} be the set of personas, \mathcal{Q} the set of 30 scored MFQ items, and R the number of repeated queries per persona–item pair. For persona p , item q , and repetition $i = 1, \dots, R$, let $y_{pqi} \in \{0, \dots, 5\}$ be the parsed rating.

For each persona–item pair we compute the sample mean and the standard deviation across repetitions

$$\bar{y}_{pq} = \frac{1}{R} \sum_{i=1}^R y_{pqi}, \quad (1)$$

$$u_{pq} = \sqrt{\frac{1}{R-1} \sum_{i=1}^R (y_{pqi} - \bar{y}_{pq})^2}, \quad (2)$$

so that u_{pq} is the standard deviation (SD) across repetitions.

Susceptibility (between-persona sensitivity) To stabilize estimates across many personas, we partition \mathcal{P} into G disjoint groups $\mathcal{P}_1, \dots, \mathcal{P}_G$ of equal size (default 10 personas per group). For each item q and group g , we compute the sample standard deviation of persona means

$$s_{gq} = \sqrt{\frac{1}{|\mathcal{P}_g| - 1} \sum_{p \in \mathcal{P}_g} (\bar{y}_{pq} - \bar{y}_{gq})^2}, \quad (3)$$

$$\bar{y}_{gq} = \frac{1}{|\mathcal{P}_g|} \sum_{p \in \mathcal{P}_g} \bar{y}_{pq}, \quad (4)$$

and average across items to obtain a group-level susceptibility sample

$$S_g = \frac{1}{|\mathcal{Q}|} \sum_{q \in \mathcal{Q}} s_{gq}. \quad (5)$$

The reported susceptibility is the mean over groups

$$S = \frac{1}{G} \sum_{g=1}^G S_g, \quad (6)$$

with its standard error estimated from the between-group variability

$$\sigma_S = \frac{\sqrt{\frac{1}{G-1} \sum_{g=1}^G (S_g - S)^2}}{\sqrt{G}}. \quad (7)$$

Foundation-specific susceptibilities reuse (4)–(7) after restricting \mathcal{Q} to the item subset \mathcal{Q}_f for foundation f .

Robustness (trial-level stability) We summarize within-pair variability by averaging the SDs in (2) over personas and items

$$\bar{u} = \frac{1}{|\mathcal{P}| |\mathcal{Q}|} \sum_{p \in \mathcal{P}} \sum_{q \in \mathcal{Q}} u_{pq}. \quad (8)$$

Our robustness index is the reciprocal

$$R = \frac{1}{\bar{u}}. \quad (9)$$

Let the (sample) standard deviation of the u_{pq} values be

$$s_u = \sqrt{\frac{1}{|\mathcal{P}| |\mathcal{Q}| - 1} \sum_{p \in \mathcal{P}} \sum_{q \in \mathcal{Q}} (u_{pq} - \bar{u})^2}. \quad (10)$$

Then the SE of \bar{u} is $\sigma_{\bar{u}} = s_u / \sqrt{|\mathcal{P}| |\mathcal{Q}|}$. Applying the delta method to (9) yields the propagated SE for robustness

$$\sigma_R = \frac{\sigma_{\bar{u}}}{\bar{u}^2}. \quad (11)$$

Foundation-level robustness repeats (8)–(11) with sums over \mathcal{Q}_f .

Cross-model normalization The z-score table (Table 1) summarizes relative performance across models. For metric $M \in \{S, R\}$ and foundation f , let $V_{mf}^{(M)}$ be model m ’s estimate with SE $\sigma_{V,mf}^{(M)}$. With across-model mean $\mu_f^{(M)}$ and SD $\sigma_f^{(M)}$, the z-score is

$$Z_{mf}^{(M)} = \frac{V_{mf}^{(M)} - \mu_f^{(M)}}{\sigma_f^{(M)}}, \quad \sigma_{Z,mf}^{(M)} = \frac{\sigma_{V,mf}^{(M)}}{\sigma_f^{(M)}}. \quad (12)$$

3 RESULTS

We present robustness and susceptibility by model and foundation, plus a z-score summary table across models.

3.1 Robustness

We quantify trial-level stability by first computing the sample standard deviation across repetitions for each persona–question pair (Eq. 2), averaging these to obtain \bar{u} (Eq. 8), and defining robustness as $R = 1/\bar{u}$ (Eq. 9) with uncertainty propagated via Eq. 11.

Observations. Robustness exhibits clear within-family structure across models. We observe a strong correlation by model family (e.g., families cluster together across foundations), with the Claude family consistently the most robust; notably, *Claude Sonnet* outperforms all others by a sizeable margin across foundations. In contrast, the Grok models are the least robust on average. We do not find a meaningful dependence on model size or SKU within families: the GPT-4.1 variants (normal, mini, nano) perform similarly, and Grok-4 versus Grok-4-fast shows no systematic size effect on robustness. These trends are visible in Figure 3 and summarized in the z-score table (Table 1).

Table 1. Overall susceptibility and robustness by model with z-scores (mean \pm SE; Z computed across models).

| Model | Robustness (\pm) | Robustness Z (\pm) | Susceptibility (\pm) | Susceptibility Z (\pm) |
|-------------------|----------------------|--------------------------|--------------------------|------------------------------|
| claude-haiku-4-5 | 32 \pm 4 | 0.2 \pm 0.1 | 0.78 \pm 0.06 | 0.6 \pm 0.8 |
| claude-sonnet-4-5 | 108 \pm 10 | 2.6 \pm 0.4 | 0.71 \pm 0.05 | −0.5 \pm 0.7 |
| gpt-4.1-mini | 11.3 \pm 0.4 | −0.41 \pm 0.01 | 0.78 \pm 0.04 | 0.5 \pm 0.7 |
| gpt-4.1-nano | 12.3 \pm 0.6 | −0.38 \pm 0.02 | 0.69 \pm 0.05 | −0.9 \pm 0.8 |
| gpt-4.1 | 14.3 \pm 0.6 | −0.32 \pm 0.02 | 0.77 \pm 0.05 | 0.4 \pm 0.7 |
| gpt-4o-mini | 12.9 \pm 0.6 | −0.36 \pm 0.02 | 0.62 \pm 0.04 | −1.9 \pm 0.5 |
| grok-4-fast | 3.33 \pm 0.06 | −0.655 \pm 0.002 | 0.86 \pm 0.05 | 1.7 \pm 0.8 |
| grok-4 | 3.31 \pm 0.05 | −0.655 \pm 0.002 | 0.75 \pm 0.03 | 0.1 \pm 0.5 |




Figure 3. Six-panel summary of robustness (inverse of average per-item standard deviation across repetitions). Top row: overall benchmark, Harm/Care, and Fairness/Reciprocity. Bottom row: In-group/Loyalty, Authority/Respect, and Purity/Sanctity. Error bars show propagated SE via delta method; higher values indicate greater rating stability.

3.2 Susceptibility

We assess between-persona sensitivity by computing within-group dispersion of persona means per item (Eq. 4), averaging across items to form group-level samples (Eq. 5), and reporting the across-group mean and its SE (Eqs. 6–7).

Observations. Susceptibility is more idiosyncratic: we do not observe strong correlation within model families, and rankings vary across foundations. The most susceptible model overall is *Grok-4-fast*, indicating larger persona-driven shifts relative to peers. Across GPT-4.1 variants there is no consistent size pattern (normal, mini, nano are comparable). See Figure 4 for the multi-foundation view and Table 1 for the corresponding z-scores.

Qualitative analysis Provide representative excerpts of reasoning (with personas anonymized) that illustrate high-susceptibility shifts versus robustly stable judgments.

4 CONCLUSION

We propose a principled benchmark for quantifying persona-driven shifts in LLM moral judgments using the MFQ-30. Our framework separates susceptibility (persona sensitivity) and robustness (rating stability), supports multiple model classes, and relies on transparent, easily repeatable procedures. Future work includes expanding persona taxonomies, stress-testing prompt formats, modeling reasoning content jointly with ratings, and correlating susceptibility with downstream alignment and safety outcomes.

REFERENCES

Abdulhai, M., Serapio-García, G., Crepy, C., Valter, D., Canny, J., and Jaques, N. Moral foundations of large language models. In Al-Onaizan, Y., Bansal, M., and Chen, Y. (eds.), *Proceedings of the 2024 Conference on Empirical Methods in Nat-*




Figure 4. Six-panel summary of moral susceptibility (mean \pm SE across persona groups). Top row: overall benchmark, Harm/Care, and Fairness/Reciprocity. Bottom row: In-group/Loyalty, Authority/Respect, and Purity/Sanctity. Higher values indicate larger persona-driven shifts in MFQ subscale scores.

ural Language Processing, pp. 17737–17752, Miami, Florida, USA, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.982. URL <https://aclanthology.org/2024.emnlp-main.982/>.

Aksoy, M. Whose morality do they speak? unraveling cultural bias in multilingual language models, 2024.

Bajpai, S., Sameer, A., and Fatima, R. Insights into moral reasoning capabilities of ai: A comparative study between humans and large language models. Research Square preprint, 2024. URL <https://doi.org/10.21203/rs.3.rs-5336157/v1>.

Costa, D. B. and Vicente, R. Deceive, detect, and disclose: Large language models play mini-mafia, 2025. URL <https://arxiv.org/abs/2509.23023>.

Ge, T., Chan, X., Wang, X., Yu, D., Mi, H., and Yu, D. Scaling synthetic data creation with 1,000,000,000 personas, 2025. URL <https://arxiv.org/abs/2406.20094>.

Graham, J., Haidt, J., and Nosek, B. A. Liberals and conservatives rely on different sets of moral foundations. *Journal of Personality and Social Psychology*, 96(5): 1029–1046, 2009. doi: 10.1037/a0015141.

Graham, J., Nosek, B. A., Haidt, J., Iyer, R., Koleva, S., and Ditto, P. H. Moral foundations questionnaire. PsycTESTS Dataset, 2011.

Haidt, J. and Graham, J. When morality opposes justice: Conservatives have moral intuitions that liberals may not

recognize. *Social Justice Research*, 20(1):98–116, 2007. doi: 10.1007/s11211-007-0034-z.

Ji, J., Chen, Y., Jin, M., Xu, W., Hua, W., and Zhang, Y. Moralbench: Moral evaluation of llms, 2025. URL <https://arxiv.org/abs/2406.04428>.

Nunes, J. L., Almeida, G. F. C. F., de Araujo, M., and Barbosa, S. D. J. Are large language models moral hypocrites? a study based on moral foundations, 2024. Final version appears in the AAAI/ACM Conference on AI, Ethics, and Society (AIES 2024).

A PROMPT FORMAT

We use a simple roleplaying wrapper with MFQ item prompts. The item texts are sourced from the Moral Foundations Questionnaire (MFQ-30; see <https://moralfoundations.org/questionnaires/>) and are not repeated here. Below we show the exact prompt formats used.

You are roleplaying as the following persona:
"<persona description>"

Please answer the following question.

When deciding whether something is right or wrong, how relevant is the following consideration to your thinking: '<CONSIDERATION>'. Please rate the consideration using this scale:
0 = not at all relevant,
1 = not very relevant,

2 = slightly relevant,
3 = somewhat relevant,
4 = very relevant,
5 = extremely relevant.

Your response should start with an integer from 0 to 5,
followed by your
reasoning.

Please indicate your level of agreement with the
following statement:
'<STATEMENT>'. Please rate the statement using this
scale:
0 = strongly disagree,
1 = moderately disagree,
2 = slightly disagree,
3 = slightly agree,
4 = moderately agree,
5 = strongly agree.

Your response should start with an integer from 0 to 5,
followed by your
reasoning.

B PERSONAS

We evaluated models under a diverse set of personas to probe persona-driven shifts in MFQ responses. We include a numbered sample below; indices match the zero-based persona identifiers (`persona_id`) used in our runs. The complete list is provided with the artifact (`personas.json`). Personas were sampled from prior work on large-scale persona generation (Ge et al., 2025).

0. A product manager focused on the integration of blockchain technology in financial services
1. A hardcore Arknights fan who is always excited to introduce new anime fans to the series
2. A marketing manager who appreciates the web developer's ability to incorporate puns into their company's website content
3. a senior tour guide specialized in Himalayan flora
4. An anthropologist exploring the cultural exchange between Viking and Irish communities through rituals and customs
5. A mission analyst who simulates and maps out the trajectories for space missions
6. A renowned world percussionist who shares their expertise and guidance
7. A Welsh aspiring screenwriter who has been following Roanne Bardsley's career for inspiration
8. The mayor of a small town who believes that the arrival of the supermarket chain will bring economic growth and job opportunities
9. A fellow book club member from a different country who has a completely different perspective on paranormal romance
10. a Slovenian industrial designer who has known Nika Zupanc since college
11. An aspiring cognitive neuroscientist seeking guidance on understanding the relationship between the brain and consciousness
12. A disabled individual who relies on the services provided by Keystone Community Resources and greatly appreciates the employee's commitment and support
13. I'm an ardent hipster music lover, DJ, and professional dancer based in New York City.
14. a hardcore fan of the Real Salt Lake soccer team
15. A self-motivated student volunteering as a research subject to contribute to the understanding of learning processes
16. A critic who argues that the author's reliance on plot twists distracts from character development
17. An inspiring fifth-grade teacher who runs the after-school cooking club
18. A high school student aspiring to become an astronaut and eagerly consumes the blogger's content for inspiration
19. an aspiring Urdu poet from India
20. A mainstream music producer who believes in sticking to industry norms and tested methods
21. A curious language enthusiast learning Latvian to better understand Baltic culture
22. A skilled tradesperson who provides vocational training in fields like construction, culinary arts, or automotive mechanics
23. A retired mass media professor staying current with marketing trends through mentorship
24. A former Miami Marlins player who played alongside Conine and formed a strong bond of camaraderie
25. A traditionalist who firmly believes Christmas should be celebrated only in December

-
26. A play-by-play announcer who excels at providing captivating player background stories during golf broadcasts
 27. A factory worker who is battling for compensation after being injured on the job due to negligence
 28. Dr. Paul R. Gregory, a Research Fellow at Stanford University's Hoover Institution, a Research Professor at the German Institute for Economic Research in Berlin, holds an endowed professorship in the Department of Economics at the University of Houston, and is emeritus chair of the International Advisory Board of the Kiev School of Economics.
 29. A science writer who relies on the geologist's knowledge and explanations for their articles
 30. A government official responsible for enforcing fair-trade regulations in the coffee industry
 31. A college professor who specializes in cognitive psychology and supports their partner's mentoring efforts
 32. A distinguished professor emeritus who has made significant contributions to the field of particle physics
 33. A filmmaker who incorporates shadow play in their movies to create a mysterious atmosphere
 34. A dedicated chef always hunting for the perfect ingredients to improve their Mediterranean cuisine recipes
 35. A young woman who is overwhelmed with the idea of planning her own wedding
 36. A fellow annoyed spouse who commiserates and shares funny anecdotes about their partners' obsessions
 37. A retired principal of a Fresh Start school in England.
 38. A talented artist who captures the fighter's journey through powerful illustrations
 39. A government official who consults the political scientist for expertise on crafting effective policy narratives
 40. a middle-aged public health official in the United States, skeptical of non-transparent practices and prefers data-led decision making
 41. A skilled jazz pianist who enjoys the challenge of interpreting gospel music
 42. A project manager who is interested in the benefits of CSS Grid and wants guidance on implementing it in future projects
 43. A political scientist writing a comprehensive analysis of global politics
 44. a fangirl who has been following Elene's career from the start.
 45. An elderly Italian man who tends to be suspicious of modern banking tools and prefers cash transactions
 46. a tech-savvy receptionist at a wellness center
 47. a resident of Torregaveta who takes local pride seriously.
 48. An experienced mobile app developer who is a minimalist.
 49. An eco-conscious local Miles from Fort Junction
 50. A current resident of the mansion whose family has a long history with the property
 51. a big fan of Ryota Muranishi who follows his games faithfully
 52. A professor specializing in cognitive neuroscience and the effects of extreme environments on the brain
 53. an ardent supporter of the different approach of politics in Greece
 54. A massage therapist exploring the connection between breathwork and relaxation techniques
 55. A retired financial professional reflecting on industry peers.
 56. A single mother who heavily relies on the mobile clinic for her family's healthcare needs and is grateful for the organizer's efforts
 57. I am a history teacher from Clare with a huge interest in local sports and cultural heritage.
 58. A marketing executive who debates about the need for less political and more lifestyle content on the blog
 59. A middle-aged aspiring novelist and music enthusiast from Edinburgh, patiently working on a draft while sipping Scottish tea on rainy afternoons.
 60. A real estate developer in Ho Chi Minh City who is always on the lookout for investment opportunities
 61. A materials scientist specializing in the development of ruggedized materials for extreme conditions
 62. A real estate agent who is always curious about the nomadic lifestyle of their relative
 63. A public policy major, focusing on healthcare disparities, inspired by their parent's work

385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439

-
64. A computer science major who often debates the impact of technology on historical data preservation
 65. An Italian local record shop owner and music enthusiast.
 66. A researcher who studies moose populations and provides insights on conservation efforts
 67. a professional iOS developer who loathes excessive typecasting
 68. A college student studying e-commerce and aids in the family business's online transition
 69. A video game developer who provides insider knowledge and references for the cosplayer's next character transformation
 70. A shy introvert discovering their voice through the art of written stories
 71. A renowned microbiologist who pioneered the field of bacterial metabolic engineering for biofuel
 72. A fresh business graduate in Pakistan
 73. A Deaf teenager struggling with their identity and navigating the hearing world
 74. A lifelong resident of Mexico City, who's elder and regularly visits Plaza Insurgentes.
 75. an ultrAslan fan, the hardcore fan group of Galatasaray SK
 76. A deeply religious family member who values their faith and seeks to share it with others
 77. An elderly retired professor who loves to learn and is interested in understanding the concept of remote work
 78. A retired historian interested in habitat laws and regulations in Texas.
 79. A film studies professor who specializes in contemporary American television and has a deep appreciation for Elmore Leonard's work.
 80. A local health clinic director seeking guidance on improving healthcare access for underserved populations
 81. A skeptical pastor from a neighboring congregation who disagrees with the preacher's teachings
 82. a Chinese retailer who sells on eBay
 83. A local real estate expert with extensive knowledge of the ancestral lands and its economic prospects

84. A prospective music student from a small town in middle America.
85. A English literature teacher trying to implement statistical analysis in grading writing assignments
86. I am a skeptical statistician who is cautious about misinterpreting results from dimensionality reduction techniques.
87. a 70-year-old veteran who served at Camp Holloway
88. A nostalgic local resident from Euxton, England who has a strong sense of community.
89. A small business owner in the beauty industry who wants to attract a specific customer base
90. A research associate who assists in analyzing retention data and identifying areas for improvement
91. A genealogist tracing the lineage of women who played influential roles during the Industrial Revolution
92. A doctoral student in development economics from Uganda
93. A mid-career Media Researcher in Ghana
94. A curriculum developer designing language courses that integrate effective pronunciation instruction
95. A dedicated music historian who helps research and uncover information about these obscure bands
96. An insurance claims adjuster who benefited from the law professor's teachings
97. A former military nurse who shares the passion for artisanal cheese and provides guidance on the business side
98. A medical professional who values personalized attention and relies on the sales representative's expertise to choose the best supplies for their practice
99. A museum curator specializing in ancient civilizations, constantly providing fascinating historical anecdotes during bridge sessions